

# Contents

<b>Advanced Machine Learning</b>	<b>1</b>
First Lecture . . . . .	1
<b>Big Data Computing</b>	<b>3</b>
First Lecture . . . . .	3
Second Lecture . . . . .	3
What's Big Data? . . . . .	3
Scaling . . . . .	3
Network Bottlenecks . . . . .	4
<b>Distributed Systems</b>	<b>5</b>
First Lecture . . . . .	5
Second Lecture . . . . .	5
Definitions . . . . .	5
Consistency . . . . .	6
Run Reconstruction . . . . .	7
<b>Models of Computation</b>	<b>11</b>
First Lesson . . . . .	11
Contact Information . . . . .	11
Course Contents . . . . .	11
Set Definition . . . . .	11
Bactus Normal Form . . . . .	12
Beta Reduction . . . . .	12
Types of Variables . . . . .	13
Extra Rules . . . . .	13



# Advanced Machine Learning

## First Lecture

A *well-posed learning problem* occurs when a computer program learns from experience  $E$ , w.r.t some task  $T$  and some performance measure  $P$ . learning occurs if its performance on  $T$ , as measured by  $P$ , improves with experience  $E$ .

Stages of intelligent systems:

1. Expert systems (1960s–1980s)
2. Unsupervised learning (1980s–1990s)
3. Data-driven learning (2000s)

This course will focus on the third stage. Increases in flexibility and model expressiveness correspond to exponential increases in the number of parameters and computational complexity, in terms of FLOPs and dataset size.



# Big Data Computing

## First Lecture

Classroom: phdlyrl

## Second Lecture

### What's Big Data?

It is sometimes used as a buzzword, but it describes an actual phenomenon.

5 V's:

1. Value: Extracting knowledge from data.
2. Volume: The amount of data. (Measuring in terabytes, petabytes, exabytes)
3. Variety: Different formats, structured (SQL tables) and Unstructured (text, images, videos).
4. Velocity: The speed at which data is generated. (usually in real-time, you can't parse it all)
5. Veracity: The quality of the data. (Is it reliable?)

### Scaling

When scaling, you have two options:

- Scaling *up*, you buy better hardware for your monolithic machine, however, moore's law is slowing down, and you will get diminishing returns, in general, you will not be able to keep up with increases in Volume/Velocity.
- Scaling *out*, you buy more machines, and distribute the workload, this is the way to go for big data, since this allows for exponential growth.

Scaling *out*, or *horizontal scaling*, also allows you for more flexibility, as once you deal with the orchestration of the workload, you can dynamically increase/reduce the number of allocated machines depending on a rolling cost/benefit analysis.

## Network Bottlenecks

A bottleneck during a request is generated when:

- There is a skew in the ratio of communication/computation
- There is a non-overlappable part of communication/computation

# Distributed Systems

## First Lecture

DA RECUPERARE !!!

A Space-Time diagram shows the evolution of a distributed system over time, it is composed of a number of timelines, each associated to a process  $p_i$ . Events, occur across this timelines, and are represented by points on the diagram. If an event  $e_1$  is causally related to another event  $e_2$ , then  $e_1$  is to the left of  $e_2$  and there is a line connecting them.

## Second Lecture

### Definitions

We now give a set of formal definitions for very natural concepts in distributed systems.

We define the *history* of a process  $p_i$  as the sequence of all events that occur in the timeline of  $p_i$ .

$$h_i = \langle e_i^1, e_i^2, \dots, e_i^n \rangle$$

We can then have the history of the computation, which is a collection of all the histories of all the processes.

$$H = \langle h_1, h_2, \dots, h_n \rangle$$

A *prefix history* is a sequence of events that occur in the timeline of a process up to a certain point.

$$h_i^k = \langle e_i^1, e_i^2, \dots, e_i^k \rangle$$

A *local state* of process  $i$  after event  $e_i^k$  is the result of the computation up to that point, and is denoted as  $\sigma_i^k$ .

A *global state* is the collection of all local states of all processes at a certain point in time.

$$\sigma^k = \langle \sigma_1^k, \sigma_2^k, \dots, \sigma_n^k \rangle$$

A *run*  $k'$  is a total re-ordering of the events such that the internal order of every process is preserved. This means that a *run* allows the interleaving of the events of different processes, but does not break the order of events in the history of a single process.

A *cut*  $C$  is a collection of the prefixes of every processes' history up to a certain point.

### Consistency

A cut  $C$  is said to be *consistent* iff:

$$e \rightarrow e' \wedge e' \in C \Rightarrow e \in C$$

Or, in other words, if an event is in the cut, then all the events that causally precede it are also in the cut.

By sequentializing the events in a run and extracting a prefix, we can trivially obtain a consistent cut.

### Consistency of Intersection

A possible course exercise is, given two cuts  $C_1$  and  $C_2$ , to prove that their intersection is also a consistent cut.

Given the definition of consistency accomodated for the intersection of two cuts:

$$e \rightarrow e' \wedge e' \in C_1 \cap C_2 \Rightarrow e \in C_1 \cap C_2$$

We know that if  $e'$  is in the intersection, then it is in both  $C_1$  and  $C_2$ . By the definition of consistency, we know that  $e$  must be in both  $C_1$  and  $C_2$ , and therefore in their intersection.

Formally:

$$e \rightarrow e' \wedge e' \in C_1 \rightarrow e \in C_1$$

$$e \rightarrow e' \wedge e' \in C_2 \rightarrow e \in C_2$$



Therefore:

$$e \rightarrow e' \wedge e' \in C_1 \cap C_2 \rightarrow e \in C_1 \wedge e \in C_2 \rightarrow e \in C_1 \cap C_2$$

### Consistency of Union

Given two consistent cuts  $C_1$  and  $C_2$ , we can prove that their union is also a consistent cut.

Given the definition of consistency accomodated for the union of two cuts:

$$e \rightarrow e' \wedge e' \in C_1 \cup C_2 \Rightarrow e \in C_1 \cup C_2$$

Then, formally:

$$e \rightarrow e' \wedge e' \in C_1 \rightarrow e \in C_1$$

$$e \rightarrow e' \wedge e' \in C_2 \rightarrow e \in C_2$$

Therefore:

$$e \rightarrow e' \wedge e' \in C_1 \cup C_2 \rightarrow e \in C_1 \vee e \in C_2 \rightarrow e \in C_1 \cup C_2$$

### Run Reconstruction

Assume that we have a set of processes  $P = \{p_1, p_2, p_3, \dots\}$ , and a monitor process  $p_0$  that is responsible for detecting deadlocks.

Everytime that an event is generated, its responsible process sends a message to the monitor, which then updates its local state.

The monitor process can then reconstruct the run by aggregating the local process states into a global state.

Assuming asynchrony and unbounded message delays, the monitor process does not reconstruct a run, since the order of the individual processes can be inverted.

If we assume a FIFO message delivery, then the monitor process can reconstruct a run. This can be trivially achieved by having a sequence number, in a system such as TCP.

This does *not* recover consistency, since the order of separate processes can still be inverted. However, this level of *local consistency* is sufficient for deadlock detection.

### Simplifying Assumptions for Consistency

Assume that I have:

1. A finite  $\delta_i$  notification delay on each process  $p_i$ .
2. A global clock  $C$  that is synchronized with all processes. We call an idealized version of this clock *Real Clock* (RC).

Then, the monitor process can offline-reconstruct a run by using the global clock to order the events.

Online-reconstruction can be achieved (and is preferred) by exploiting the additional  $\delta_i$  information. Once we receive a notification, we await time  $\max \delta_i$  before committing the event to the global state, in the meantime, we store the event in a buffer that is ordered by the global clock.

### Real Clock Property

The clock property relates the *happens before* relation of messages to the real time of the system.

$$e \rightarrow e' \Rightarrow TS(e) < TS(e')$$

We can obtain a trivial clock by designing a system as such:

1. Each process has associates a sequence number  $k \in \mathbb{N}$  to each event.
2. Each event  $i$ , which receives a message  $m$  from event  $j$ , has a sequence number given by the formula  $\max(k_i, k_j) + 1$ .

This gives us the clock property *by construction*, since each event has a sequence number that is always greater than the sequence number of the event that caused it and of any event that it is casually related to.

This device is called a *Lamport Clock*, or a *Logical Clock*.

### History of an event

Given an event  $e$ , we can define its history as the sequence of events that causally precede it. These are all the events that the monitor process must have received before the event  $e$  is committed to the global state.

$$\Theta(e) = \{e' \mid e' \rightarrow e\}$$

This is a *consistent* cut.

### Opaque detection of consistent cuts

Given a notification about an event, for example,  $e_5^4$ , the monitor process can ensure that  $e_1^4$  can be committed to global state by checking that for every process, we have received a notification with lamport clock  $\geq 4$ .

We might have some processes that consistently lag behind, and therefore do not reach the lamport clock in a timely manner. We can trivially solve this by sending a liveness check to all processes for which a sufficiently high lamport number has not been received. If the process is alive but not working, it will re-adjust its lamport clock and send a notification.

This will not impact throughput, since the monitor process is not a bottleneck, and the liveness check is a very lightweight operation, that is performed only when necessary.

### Building a Better Clock

We can build a *better clock*, that is, one with the enhanced property:

$$e \rightarrow e' \iff TS(e) < TS(e')$$

This is called a *Vector Clock*.

It works by communicating the history of the current event in a compact way, we do this by sending a vector of length  $n$  where  $n$  is the number of processes in the system. The vector's entries are updated to be the highest lamport number detected for that relative process, either from a message that has been received, or from the local lamport clock (in the case of the current process).

This allows us to check whether:

$$VC(e) \subseteq VC(e')$$

$$e \rightarrow e' \iff VC(e) \subseteq VC(e')$$

When messages are sent, the vector clock is attached to the message, and the receiving process updates its vector clock by taking the maximum of the two vectors, and summing 1 to the current process.

The monitor process  $p_0$  keeps track of its own vector, where each entry is the highest lamport number that it has received from each process. When it receives a notification, if the vector clock is consistent with the monitor's vector clock, then the event is committed to the global state.

Since vectors are *not* sets, we have to use component-wise comparison to check for causality.

$$VC(e) < VC(e') \iff \forall i \in \{1, 2, \dots, n\} \quad VC(e)_i < VC(e')_i$$

So we are back to the original definition of the happens before relation.

# Models of Computation

## First Lesson

### Contact Information

Prof email: [piperno@di.uniroma1.it](mailto:piperno@di.uniroma1.it)

*Actual* lecture times:

- Wednesday 13:30 - 15:00
- Thursday 16:00 - 17:30, or 16:15 - 17:55

### Course Contents

The course content will focus on **functional programming** and  $\lambda$ -calculus.

The main application of these languages is that of *function application*, these languages are devoid of *assignment* semantics, and are therefore called *pure*.

The main ingredients of our programs will be:

- Variables, which are usually denoted by lowercase letters
- Starting from these, you obtain the set of all programs, which are called  $\lambda$ -terms, this set is denoted as  $\Lambda$ .

### Set Definition

We can provide an inductive definition of a set of  $\lambda$ -terms,  $\Lambda$ :

$$\frac{x \in V}{x \in \Lambda} \quad (\text{var})$$

$$\frac{M \in \Lambda \quad N \in \Lambda}{(MN) \in \Lambda} \quad (\text{app})$$

$$\frac{M \in \Lambda \quad x \in V}{\lambda x.M \in \Lambda} \quad (\text{abs})$$

Applying these inductive rules (variables, application, abstraction).

### Bactus Normal Form

Another way of describing a set of lambda terms is by using Bactus Normal Form (BNF).

$$\Lambda :: Var | \Lambda \Lambda | \lambda Var. \Lambda$$

Which is essentially a grammar for the set of lambda terms.

Lambda calculus is left-associative:

$$((xy)z) = xyz \neq x(yz)$$

All functions are unary (we assume currying). For example, the function  $f(x, y)$  is represented as:

$$\lambda x. \lambda y. fxy$$

.

Functions in lambda calculus can be applied to other functions or themselves, they can also return functions.

### Beta Reduction

The main operation in lambda calculus is *beta reduction*, which is the application of a function to an argument.

$$\frac{(\lambda x. M)N}{redex} \rightarrow_{\beta} M[N/x]$$

Where  $M[N/x]$  is the result of substituting all occurrences of  $x$  in  $M$  with  $N$ .

Some other examples:

$$(\lambda x. x)y \rightarrow_{\beta} y$$

$$(\lambda x. xx)y \rightarrow_{\beta} yy$$

$$(\lambda xy. yx)(\lambda u. u) \rightarrow_{\beta} \lambda y. y(\lambda u. u)$$

$$(\lambda xy. yx)(\lambda t. y) \rightarrow_{\beta} \lambda y. y(\lambda t. y)$$

This rule can be applied in any context in which it appears.

## Types of Variables

We distinguish two kinds of variables:

- Free variables: variables that are not bound by an abstraction
- Bound variables: variables that are bound by an abstraction

For example, in the term  $\lambda x.xy$ ,  $y$  is a free variable, while  $x$  is a bound variable.

Bound variables can be renamed, whereas for free variables the naming is relevant.

$$\begin{cases} FV(x) = \{x\} \\ FV(MN) = FV(M) \cup FV(N) \\ FV(\lambda x.M) = FV(M) - \{x\} \end{cases}$$

A set of lambda terms where  $LM(\Lambda) = \emptyset$  is called *closed*.

## Extra Rules

$$\begin{aligned} (\mu) \quad & \frac{M \rightarrow_{\beta} M'}{NM \rightarrow_{\beta} NM'} \\ (\nu) \quad & \frac{M \rightarrow_{\beta} M'}{MN \rightarrow_{\beta} M'N} \\ (\xi) \quad & \frac{M \rightarrow M'}{\lambda x.M \rightarrow \lambda x.M'} \end{aligned}$$

These rules allow us to select redexes in a context-free manner in the middle of our lambda term. We can then choose the order of evaluation of our redexes, while still taking care of the left-associative order of precedence. Our calculus is therefore not *determinate* but is still *deterministic*, meaning that there may be multiple reduction strategies but they all lead to the same result.

This corollary is called the *Church Rosser Theorem*, discovered in 1936.

In general, a call-by-value-like semantic is preferable when choosing evaluation paths, as it clears the most amount of terms as early as possible.

- Call By Value is *efficient*
- Call By Name is *complete*, if the lambda term is normalizable

